

PROBABILISTIC MACHINE LEARNING

LECTURE 23

TUNING INFERENCE ALGORITHMS


Philipp Hennig

07 July 2020

EBERHARD KARLS
UNIVERSITÄT
TÜBINGEN



FACULTY OF SCIENCE
DEPARTMENT OF COMPUTER SCIENCE
CHAIR FOR THE METHODS OF MACHINE LEARNING



| # | date | content | Ex | # | date | content | Ex |
|----|--------|------------------------------|----|----|--------|-----------------------------|----|
| 1 | 20.04. | Introduction | 1 | 14 | 09.06. | Generalized Linear Models | |
| 2 | 21.04. | Reasoning under Uncertainty | | 15 | 15.06. | Exponential Families | 8 |
| 3 | 27.04. | Continuous Variables | 2 | 16 | 16.06. | Graphical Models | |
| 4 | 28.04. | Monte Carlo | | 17 | 22.06. | Factor Graphs | 9 |
| 5 | 04.05. | Markov Chain Monte Carlo | 3 | 18 | 23.06. | The Sum-Product Algorithm | |
| 6 | 05.05. | Gaussian Distributions | | 19 | 29.06. | Example: Modelling Topics | 10 |
| 7 | 11.05. | Parametric Regression | 4 | 20 | 30.06. | Mixture Models | |
| 8 | 12.05. | Learning Representations | | 21 | 06.07. | EM | 11 |
| 9 | 18.05. | Gaussian Processes | 5 | 22 | 07.07. | Variational Inference | |
| 10 | 19.05. | Understanding Kernels | | 23 | 13.07. | Tuning Inference Algorithms | 12 |
| 11 | 26.05. | Gauss-Markov Models | | 24 | 14.07. | Kernel Topic Models | |
| 12 | 25.05. | An Example for GP Regression | 6 | 25 | 20.07. | Outlook | |
| 13 | 08.06. | GP Classification | 7 | 26 | 21.07. | Revision | |



Designing a probabilistic machine learning method:

1. get the **data**
 - 1.1 try to collect as much meta-data as possible
2. build the **model**
 - 2.1 identify quantities and datastructures; assign names
 - 2.2 design a generative process (graphical model)
 - 2.3 assign (conditional) distributions to factors/arrows (use exponential families!)
3. design the **algorithm**
 - 3.1 consider conditional independence
 - 3.2 try standard methods for early experiments
 - 3.3 run unit-tests and sanity-checks
 - 3.4 identify bottlenecks, find customized approximations and refinements

Framework:

$$\int p(x_1, x_2) dx_2 = p(x_1)$$

$$p(x_1, x_2) = p(x_1 | x_2)p(x_2)$$

$$p(x | y) = \frac{p(y | x)p(x)}{p(y)}$$

Modelling:

- ▶ graphical models
- ▶ Gaussian distributions
- ▶ (deep) learnt representations
- ▶ Kernels
- ▶ Markov Chains
- ▶ Exponential Families / Conjugate Priors
- ▶ Factor Graphs & Message Passing

Computation:

- ▶ Monte Carlo
- ▶ Linear algebra / Gaussian inference
- ▶ maximum likelihood / MAP
- ▶ Laplace approximations
- ▶ EM / variational approximations






The Data and an idea for a model



Making Assumptions

Our Data, our model




- ▶ a corpus of D documents
- ▶ each containing I_d words from a vocabulary of V words
- ▶ assumed to consist of K topics



The Model





To draw l_d words $w_{di} \in [1, \dots, V]$ of document $d \in [1, \dots, D]$:

- ▶ Draw K topic distributions θ_k over V words from
- ▶ Draw D document distributions over K topics from
- ▶ Draw topic assignments c_{dik} of word w_{di} from
- ▶ Draw word w_{di} from


Useful notation: $n_{dkv} = \#\{i : w_{di} = v, c_{dik} = 1\}$. Write $n_{dk\cdot} := [n_{dk1}, \dots, n_{dkV}]$ and $n_{dk\cdot} = \sum_v n_{dkv}$, etc.

$$p(\Theta \mid \beta) = \prod_{k=1}^K \mathcal{D}(\theta_k; \beta_k)$$


$$p(\Pi \mid \alpha) = \prod_{d=1}^D \mathcal{D}(\pi_d; \alpha_d)$$

$$p(C \mid \Pi) = \prod_{i,d,k} \pi_{dk}^{c_{dik}}$$

$$p(w_{di} = v \mid c_{di}, \Theta) = \prod_k \theta_{kv}^{c_{dik}}$$



$$\begin{aligned}
 p(C, \Pi, \Theta, W) &= \underbrace{\left(\prod_{d=1}^D \mathcal{D}(\pi_d; \alpha_d) \right)}_{p(\Pi|\alpha)} \cdot \underbrace{\left(\prod_{d=1}^D \prod_{i=1}^{I_d} \left(\prod_{k=1}^K \pi_{dk}^{C_{dik}} \right) \right)}_{p(C|\Pi)} \cdot \underbrace{\left(\prod_{d=1}^D \prod_{i=1}^{I_d} \left(\prod_{k=1}^K \theta_{kw_{di}}^{C_{dik}} \right) \right)}_{p(W|C, \Theta)} \cdot \underbrace{\left(\prod_{k=1}^K \mathcal{D}(\theta_k; \beta_k) \right)}_{p(\Theta|\beta)} \\
 &= \underbrace{\left(\prod_{d=1}^D \mathcal{D}(\pi_d; \alpha_d) \right)}_{p(\Pi|\alpha)} \cdot \underbrace{\left(\prod_{d=1}^D \prod_{i=1}^{I_d} \left(\prod_{k=1}^K (\pi_{dk} \theta_{kw_{di}})^{C_{dik}} \right) \right)}_{p(W, C|\Theta, \Pi)} \cdot \underbrace{\left(\prod_{k=1}^K \mathcal{D}(\theta_k; \beta_k) \right)}_{p(\Theta|\beta)} \\
 &= \left(\prod_{d=1}^D \frac{\Gamma(\sum_k \alpha_{dk})}{\prod_k \Gamma(\alpha_{dk})} \prod_{k=1}^K \pi_{dk}^{\alpha_{dk}-1+n_{dk}} \right) \cdot \left(\prod_{k=1}^K \frac{\Gamma(\sum_v \beta_{kv})}{\prod_v \Gamma(\beta_{kv})} \prod_{v=1}^V \theta_{kv}^{\beta_{kv}-1+n_{kv}} \right)
 \end{aligned}$$




$$p(C, \Pi, \Theta, W) = \left(\prod_{d=1}^D \mathcal{D}(\pi_d; \alpha_d) \right) \cdot \left(\prod_{d=1}^D \prod_{i=1}^{I_d} \left(\prod_{k=1}^K \pi_{dk}^{c_{dik}} \right) \right) \cdot \left(\prod_{d=1}^D \prod_{i=1}^{I_d} \left(\prod_{k=1}^K \theta_{kw_{di}}^{c_{dik}} \right) \right) \cdot \left(\prod_{k=1}^K \mathcal{D}(\theta_k; \beta_k) \right)$$

- If we had Π, Θ (which we don't), then the posterior $p(C \mid \Theta, \Pi, W)$ would be easy:

$$p(C \mid \Theta, \Pi, W) = \frac{p(W, C, \Theta, \Pi)}{\sum_C p(W, C, \Theta, \Pi)} = \prod_{d=1}^D \prod_{i=1}^{I_d} \frac{\prod_{k=1}^K (\pi_{dk} \theta_{kw_{di}})^{c_{dik}}}{\sum_{k'} (\pi_{dk'} \theta_{k'w_{di}})^{c_{dik}}}$$

- note that this conditional independence can easily be read off from the above graph!



$$p(C, \Pi, \Theta, W) = \left(\prod_{d=1}^D \frac{\Gamma(\sum_k \alpha_{dk})}{\prod_k \Gamma(\alpha_{dk})} \prod_{k=1}^K \pi_{dk}^{\alpha_{dk}-1+n_{dk}} \right) \cdot \left(\prod_{k=1}^K \frac{\Gamma(\sum_v \beta_{kv})}{\prod_v \Gamma(\beta_{kv})} \prod_{v=1}^V \theta_{kv}^{\beta_{kv}-1+n_{kv}} \right)$$

- If we had C (which we don't), then the posterior $p(\Theta, \Pi \mid C, W)$ would be easy:

$$\begin{aligned} p(\Theta, \Pi \mid C, W) &= \frac{p(C, W, \Pi, \Theta)}{\int p(\Theta, \Pi, C, W) d\Theta d\Pi} = \frac{(\prod_d \mathcal{D}(\pi_d; \alpha_d) (\prod_k \pi_{dk}^{n_{dk}})) (\prod_k \mathcal{D}(\theta_k; \beta_k) (\prod_v \theta_{kv}^{n_{kv}}))}{p(C, W)} \\ &= \left(\prod_d \mathcal{D}(\pi_d; \alpha_{d:} + n_{d:}) \right) \left(\prod_k \mathcal{D}(\theta_k; \beta_{k:} + n_{k:}) \right) \end{aligned}$$

- note that this conditional independence **can not** easily be read off from the above graph!



The Algorithms



Iterate between (recall $n_{dkv} = \#\{i : w_{di} = v, c_{ijk} = 1\}$)

$$\Theta \sim p(\Theta \mid \mathcal{C}, W) = \prod_k \mathcal{D}(\theta_k; \beta_k + n_{\cdot k})$$

$$\Pi \sim p(\Pi \mid \mathcal{C}, W) = \prod_d \mathcal{D}(\pi_d; \alpha_d + n_{d \cdot})$$

$$\mathcal{C} \sim p(\mathcal{C} \mid \Theta, \Pi, W) = \prod_{d=1}^D \prod_{i=1}^{l_d} \frac{\prod_{k=1}^K (\pi_{dk} \theta_{kw_{di}})^{c_{dik}}}{\sum_{k'} (\pi_{dk'} \theta_{k'w_{di}})}$$

- ▶ This is *comparably* easy to implement because there are libraries for sampling from Dirichlet's, and discrete sampling is trivial. All we have to keep around are the counts n (which are sparse!) and Θ, Π (which are comparably small). Thanks to factorization, much can also be done in parallel!
- ▶ Unfortunately, this sampling scheme is relatively slow to move out of initialization, because z depends strongly on θ, π and vice versa.
- ▶ properly vectorizing the code is important for speed

To maximize $p(\Theta, \Pi \mid W)$, consider (where $\gamma_{dik} := \pi_{dk} \theta_{kw_{di}}$ and $\tilde{\gamma}_{dik} := \gamma_{dik} / \sum_{k'} \gamma_{dik'}$)

$$p(C \mid \Theta, \Pi, W) = \prod_{d=1}^D \prod_{i=1}^{I_d} \frac{\prod_{k=1}^K (\pi_{dk} \theta_{kw_{di}})^{c_{dik}}}{\sum_{k'} (\pi_{dk'} \theta_{k'w_{di}})} = \prod_{d=1}^D \prod_{i=1}^{I_d} \prod_{k=1}^K \tilde{\gamma}_{dik}^{c_{dik}}$$

And **Maximize** the Expected complete log posterior

$$\mathbb{E}_{p(C|\gamma)}[\log p(\Theta, \Pi \mid W, C)] = \sum_d \sum_k (\tilde{\gamma}_{d \cdot k} + \alpha_{dk}) \cdot \log \pi_{dk} + \sum_k \sum_v \left(\beta_{kv} + \sum_d \sum_i \mathbb{I}(w_{di} = v) \tilde{\gamma}_{dik} \right) \log \theta_{kv}$$

at

$$\pi_{dk} = \frac{\tilde{\gamma}_{d \cdot k} + \alpha_{dk}}{\tilde{\gamma}_{d \cdot \cdot} + \alpha_{d \cdot}} \quad \text{and} \quad \theta_{kv} = \frac{\beta_{kv} + \sum_{d,i} \mathbb{I}(w_{di} = v) \tilde{\gamma}_{dik}}{\beta_{k \cdot} + \sum_{v'} \sum_{d,i} \mathbb{I}(w_{di} = v') \tilde{\gamma}_{dik}}$$

and repeat.

To find the distribution q the minimizes, subject to $q(\Pi, \Theta, C) = q(\Pi, \Theta) \cdot q(C)$

$$D_{\text{KL}}(q(\Pi, \Theta, C) \| p(\Pi, \Theta, C | W)) = \int q \log \frac{q}{p} dq$$

or, equivalently, maximizes the ELBO

$$\mathcal{L}(q) = \int q(C, \Pi, \Theta) \log \left(\frac{p(C, \Pi, \Theta, W)}{q(C, \Pi, \Theta)} \right) dC d\Pi d\Theta$$

set $\log q(C) = \mathbb{E}_{q(\Pi, \Theta)}(\log p(C, \Pi, \Theta, W))$ and vice versa, to get...

$$q(\boldsymbol{\pi}_d) = \mathcal{D} \left(\boldsymbol{\pi}_d; \tilde{\alpha}_{dk} := \left[\alpha_{dk} + \sum_{i=1}^{I_d} \tilde{\gamma}_{dik} \right]_{k=1, \dots, K} \right) \quad \forall d = 1, \dots, D$$

$$q(\boldsymbol{\theta}_k) = \mathcal{D} \left(\boldsymbol{\theta}_k; \tilde{\beta}_{kv} := \left[\beta_{kv} + \sum_d^D \sum_{i=1}^{I_d} \tilde{\gamma}_{dik} \mathbb{I}(w_{di} = v) \right]_{v=1, \dots, V} \right) \quad \forall k = 1, \dots, K$$

$$q(\mathbf{c}_{di}) = \prod_k \tilde{\gamma}_{dik}^{c_{dik}}, \quad \forall d \quad i = 1, \dots, I_d$$

where $\tilde{\gamma}_{dik} = \gamma_{dik} / \sum_k \gamma_{dik}$ and (note that $\sum_k \tilde{\alpha}_{dk} = \text{const.}$)

$$\gamma_{dik} = \exp \left(\mathbb{E}_{q(\boldsymbol{\pi}_{dk})} (\log \pi_{dk}) + \mathbb{E}_{q(\boldsymbol{\theta}_{di})} (\log \theta_{kw_{di}}) \right) = \exp \left(F(\tilde{\alpha}_{jk}) + F(\tilde{\beta}_{kw_{di}}) - F \left(\sum_v \tilde{\beta}_{kv} \right) \right)$$

and repeat



An idea.



- ▶ Consider the exponential family $p_w(x | w) = \exp [\phi(x)^\top w - \log Z(w)]$
- ▶ its conjugate prior is the exponential family $F(\alpha, \nu) = \int \exp(\alpha^\top w - \nu^\top \log Z(w)) dw$

$$p_\alpha(w | \alpha, \nu) = \exp \left[\begin{pmatrix} w \\ -\log Z(w) \end{pmatrix}^\top \begin{pmatrix} \alpha \\ \nu \end{pmatrix} - \log F(\alpha, \nu) \right]$$

$$\text{because } p_\alpha(w | \alpha, \nu) \prod_{i=1}^n p_w(x_i | w) \propto p_\alpha \left(w \middle| \alpha + \sum_i \phi(x_i), \nu + n \right)$$

- ▶ and the predictive is

$$\begin{aligned}
 p(x) &= \int p_w(x | w) p_\alpha(w | \alpha, \nu) dw = \int e^{(\phi(x) + \alpha)^\top w - (\nu + 1) \log Z(w) - \log F(\alpha, \nu)} dw \\
 &= \frac{F(\phi(x) + \alpha, \nu + 1)}{F(\alpha, \nu)}
 \end{aligned}$$

Exponential Families, among other things (see also last lecture) provide **conjugate priors** for standard distributions (Lectures 2,15)

- ▶ Consider the exponential family $p(c \mid \pi) = \exp \left[c^\top (\log \pi) - \log \sum_k \pi_k \right]$
- ▶ its conjugate prior is the exponential family $B(\alpha) = \int \exp(\alpha^\top \log \pi - \nu \cdot 0) d\pi$

$$\mathcal{D}(\pi \mid \alpha) = \exp [\log \pi^\top \alpha - \log B(\alpha)]$$

$$\text{because } \mathcal{D}(\pi \mid \alpha) \prod_{i=1}^n \pi^{c_i} \propto \mathcal{D} \left(\pi \mid \alpha + \sum_i c_i \right)$$

- ▶ and the predictive is

$$p(c) = \int p(c \mid \pi) \mathcal{D}(\pi \mid \alpha) d\pi = \int e^{(c+\alpha)^\top (\log \pi) + \log B(\alpha)} d\pi = \frac{B(c + \alpha)}{B(\alpha)}$$

Exponential Families, among other things (see also last lecture) provide **conjugate priors** for standard distributions (Lectures 2,15)



Recall $\Gamma(x + 1) = x \cdot \Gamma(x) \ \forall x \in \mathbb{R}_+$

$$\begin{aligned}
 p(C, \Pi, \Theta, W) &= \left(\prod_{d=1}^D \frac{\Gamma(\sum_k \alpha_{dk})}{\prod_k \Gamma(\alpha_{dk})} \prod_{k=1}^K \pi_{dk}^{\alpha_{dk}-1+n_{dk}} \right) \cdot \left(\prod_{k=1}^K \frac{\Gamma(\sum_v \beta_{kv})}{\prod_v \Gamma(\beta_{kv})} \prod_{v=1}^V \theta_{kv}^{\beta_{kv}-1+n_{kv}} \right) \\
 &= \left(\prod_{d=1}^D \frac{B(\alpha_d + n_{d\cdot})}{B(\alpha_d)} \mathcal{D}(\pi_d; \alpha_d + n_{d\cdot}) \right) \cdot \left(\prod_{k=1}^K \frac{B(\beta_k + n_{\cdot k})}{B(\beta_k)} \mathcal{D}(\theta_k; \beta_k + n_{\cdot k}) \right)
 \end{aligned}$$

$$\begin{aligned}
 p(C, W) &= \left(\prod_{d=1}^D \frac{B(\alpha_d + n_{d\cdot})}{B(\alpha_d)} \right) \cdot \left(\prod_{k=1}^K \frac{B(\beta_k + n_{\cdot k})}{B(\beta_k)} \right) \\
 &= \left(\prod_d \frac{\Gamma(\sum_{k'} \alpha_{dk'})}{\Gamma(\sum_{k'} \alpha_{dk'} + n_{dk'})} \prod_k \frac{\Gamma(\alpha_{dk} + n_{dk})}{\Gamma(\alpha_{dk})} \right) \left(\prod_k \frac{\Gamma(\sum_v \beta_{kv})}{\Gamma(\sum_v \beta_{kv} + n_{kv})} \prod_v \frac{\Gamma(\beta_{kv} + n_{kv})}{\Gamma(\beta_{kv})} \right)
 \end{aligned}$$

$$p(c_{dik} = 1 \mid C^{\setminus di}, W) = \frac{(\alpha_{dk} + n_{dk}^{\setminus di})(\beta_{kw_{di}} + n_{kw_{di}}^{\setminus di})(\sum_v \beta_{kv} + n_{kv}^{\setminus di})^{-1}}{\sum_{k'} (\alpha_{dk'} + n_{dk'}^{\setminus di}) \cdot \sum_{w'} (\beta_{kw'} + n_{kw'}^{\setminus di}) \cdot \sum_{v'} (\beta_{kv'} + n_{kv'}^{\setminus di})^{-1}}$$

A Collapsed Gibbs Sampler for LDA

It pays off to look closely at the math!

T. L. Griffiths & M. Steyvers, *Finding scientific topics*, PNAS 101/1 (4/2004), 5228–5235

$$p(C, W) = \left(\prod_d \frac{\Gamma(\sum_k \alpha_{dk})}{\Gamma(\sum_k \alpha_{dk} + n_{d\cdot})} \prod_k \frac{\Gamma(\alpha_{dk} + n_{dk\cdot})}{\Gamma(\alpha_{dk})} \right) \left(\prod_k \frac{\Gamma(\sum_v \beta_{kv})}{\Gamma(\sum_v \beta_{kv} + n_{\cdot kv})} \prod_v \frac{\Gamma(\beta_{kv} + n_{\cdot kv})}{\Gamma(\beta_{kv})} \right)$$

A **collapsed** sampling method can converge much faster by eliminating the latent variables that mediate between individual data.

```
1 procedure LDA( $W, \alpha, \beta$ )  
2    $\gamma_{dkv} \leftarrow 0 \ \forall d, k, v$  // initialize counts  
3   while true do  
4     for  $d = 1, \dots, D; i = 1, \dots, I_d$  do // can be parallelized  
5        $c_{di} \propto (\alpha_{dk} + n_{dk\cdot}^{di})(\beta_{kw_{di}} + n_{\cdot kw_{di}}^{di})(\sum_v \beta_{kv} + n_{\cdot kv}^{di})^{-1}$  // sample assignment  
6        $n \leftarrow \text{UPDATECOUNTS}(c_{di})$  // update counts (check whether first pass or repeat)  
7     end for  
8   end while  
9 end procedure
```

Collapsed Sampling is quite efficient

The Mean Field argument

[figure: T. L. Griffiths & M. Steyvers, *Finding scientific topics*, PNAS **101**/1 (4/2004), 5228–5235]



Thomas Griffiths

image: Princeton U



Mark Steyvers

image: UC Irvine

The collapsed sampler operates on the **mean field**

$$p(C | W) = \int p(C | \Theta, \Pi, W) p(\Theta, \Pi | W) d\Theta d\Pi$$

The *expected* value of the variables Θ, Π that mediate between the “particles” (words). This works well because each word’s topic is approximately independent of all individual other words’ topics (but together they create the whole thing).

- Deriving our variational bound, we previously imposed the factorization

$$q(\Pi, \Theta, C) = q(\Pi, \Theta) \cdot q(C), \quad \text{but can we get away with less? Like,}$$
$$q(\Pi, \Theta, C) = q(\Theta, \Pi | C) \cdot q(C)$$

- Note $p(C, \Theta, \Pi | W) = p(\Theta, \Pi | C, W)p(C | W)$. So when we minimize

$$\begin{aligned} D_{\text{KL}}(q(\Pi, \Theta, C) \| p(\Pi, \Theta, C | W)) &= \int q(\Pi, \Theta | C) q(C) \log \left(\frac{q(\Pi, \Theta | C) q(C)}{p(\Pi, \Theta | C, W) p(C | W)} \right) dC d\Pi d\Theta \\ &= \int q(\Pi, \Theta | C) q(C) \left[\log \left(\frac{q(\Pi, \Theta | C)}{p(\Pi, \Theta | C, W)} \right) + \log \left(\frac{q(C)}{p(C | W)} \right) \right] dC d\Pi d\Theta \\ &= D_{\text{KL}}(q(\Pi, \Theta | C) \| p(\Pi, \Theta | C, W)) + D_{\text{KL}}(q(C) \| p(C | W)) \end{aligned}$$

we will just get $q(\Theta, \Pi) = p(\Theta, \Pi | C, W)$ and the bound will be *tight* in Π, Θ .

$$p(C, W) = \left(\prod_d \frac{\Gamma(\sum_k \alpha_{dk})}{\Gamma(\sum_k \alpha_{dk} + n_{dk\cdot})} \prod_k \frac{\Gamma(\alpha_{dk} + n_{dk\cdot})}{\Gamma(\alpha_{dk})} \right) \left(\prod_k \frac{\Gamma(\sum_v \beta_{kv})}{\Gamma(\sum_v \beta_{kv} + n_{\cdot kv})} \prod_v \frac{\Gamma(\beta_{kv} + n_{\cdot kv})}{\Gamma(\beta_{kv})} \right)$$

- The remaining **collapsed variational bound** (ELBO) becomes

$$\mathcal{L}(q) = \int q(C) \log p(C, W) dC + \mathbb{H}(q(C))$$

- because we make strictly less assumptions about q than before, we will get a strictly better approximation to the true posterior!
- The bound is maximized for c_{di} if

$$\log q(c_{di}) = \mathbb{E}_{q(C \setminus \{di\})}(\log p(C, W)) + \text{const.}$$

Why didn't we do this earlier?

- Note that $c_{di} \in \{0; 1\}^K$ and $\sum_k c_{dik} = 1$. So $q(c_{di}) = \prod_k \gamma_{dik}$ with $\sum_k \gamma_{dik} = 1$
- Also: $\Gamma(\alpha + n) = \prod_{\ell=0}^{n-1} (\alpha + \ell)$, thus $\log \Gamma(\alpha + n) = \sum_{\ell=0}^{n-1} \log(\alpha + \ell)$

$$p(C, W) = \left(\prod_d \frac{\Gamma(\sum_k \alpha_{dk})}{\Gamma(\sum_k \alpha_{dk} + n_{dk.})} \prod_k \frac{\Gamma(\alpha_{dk} + n_{dk.})}{\Gamma(\alpha_{dk})} \right) \left(\prod_k \frac{\Gamma(\sum_v \beta_{kv})}{\Gamma(\sum_v \beta_{kv} + n_{.kv})} \prod_v \frac{\Gamma(\beta_{kv} + n_{.kv})}{\Gamma(\beta_{kv})} \right)$$

$$\log q(c_{di}) = \mathbb{E}_{q(C \setminus di)} (\log p(C, W)) + \text{const.}$$

$$\log \gamma_{dik} = \log q(c_{dik} = 1)$$

$$= \mathbb{E}_{q(C \setminus di)} \left[\log \Gamma(\alpha_{dk} + n_{dk.}) + \log \Gamma(\beta_{kw_{di}} + n_{.kw_{di}}) - \log \Gamma \left(\sum_v \beta_{kv} + n_{.kv} \right) \right] + \text{const.}$$

$$= \mathbb{E}_{q(C \setminus di)} \left[\log(\alpha_{dk} + n_{dk.}^{di}) + \log(\beta_{kw_{di}} + n_{.kw_{di}}^{di}) - \log \left(\sum_v \beta_{kv} + n_{.kv}^{di} \right) \right] + \text{const.}$$

(note all terms in $p(C, W)$ that don't involve c_{dik} can be moved into the constant, as can all sums over k .)

We can also *add* terms to const., such as $\sum_{\ell=0}^{n_{dk.}^{di}-1} \log(\alpha + \ell)$, effectively cancelling terms in $\log \Gamma$)



$$\gamma_{dik} \propto \exp \left(\mathbb{E}_{q(C \setminus di)} \left[\log(\alpha_{dk} + n_{dk\cdot}^{di}) + \log(\beta_{kw_{di}} + n_{\cdot kw_{di}}^{di}) - \log \left(\sum_v \beta_{kv} + n_{\cdot kv}^{di} \right) \right] \right)$$

- Under $q(C) = \prod_{di} c_{di}$, the counts $n_{dk\cdot}$ are sums of independent Bernoulli variables (i.e. they have a **multinomial** distribution). Computing their expected logarithm is tricky ($\mathcal{O}(n_{d\cdot}^2)$):

$$\mathbb{H}(q(n_{dk\cdot})) = \mathbb{E}[\log n_{dk\cdot}] = -\log(l_d!) - l_d \sum_k \gamma_{dk\cdot} \log(\gamma_{dk\cdot}) + \sum_{k=1}^K \sum_{n_{dk\cdot}=1}^{l_d} \binom{l_d}{n_{dk\cdot}} \gamma_{dk\cdot}^{n_{dk\cdot}} (1 - \gamma_{dk\cdot})^{l_d - n_{dk\cdot}} \log(n_{dk\cdot}!)$$

- That's likely why the original paper (and `scikit-learn`) don't do this.

If arithmetic doesn't work, try creativity!



Yee Whye Teh, David Newman & Max Welling, NeurIPS 2007



Yee Whye Teh

image: Oxford U



Max Welling


image: U v Amsterdam

$$\gamma_{dik} \propto \exp \left(\mathbb{E}_{q(C \setminus di)} \left[\log(\alpha_{dk} + n_{dk \cdot}^{di}) + \log(\beta_{kw_{di}} + n_{\cdot kw_{di}}^{di}) - \log \left(\sum_v \beta_{kv} + n_{\cdot kv}^{di} \right) \right] \right)$$



The probability measure of $R = \sum_i^N x_i$ with discrete x_i of probability f is

$$\begin{aligned} P(R = r \mid f, N) &= \frac{N!}{(N-r)! \cdot r!} \cdot f^r \cdot (1-f)^{N-r} \\ &= \binom{N}{r} \cdot f^r \cdot (1-f)^{N-r} \\ &\approx \mathcal{N}(r; Nr, Nr(1-f)) \end{aligned}$$



If arithmetic doesn't work, try creativity!



Yee Whye Teh, David Newman & Max Welling, NeurIPS 2007



Yee Whye Teh

image: Oxford U



Max Welling

image: U v Amsterdam

but the CLT applies! So a Gaussian approximation should be good:

$$p(n_{dk.}^{di}) \approx \mathcal{N}(n_{dk.}^{di}; \mathbb{E}_q[n_{dk.}^{di}], \text{var}_q[n_{dk.}^{di}]) \quad \text{with} \quad \mathbb{E}_q[n_{dk.}^{di}] = \sum_{j \neq i} \gamma_{dkj}, \quad \text{var}_q[n_{dk.}^{di}] = \sum_{j \neq i} \gamma_{dkj}(1 - \gamma_{dkj})$$

If arithmetic doesn't work, try creativity!



Yee Whye Teh, David Newman & Max Welling, NeurIPS 2007



Yee Whye Teh

image: Oxford U



Max Welling

image: U v Amsterdam

$$\mathbb{E}_q[\log(\alpha_{dk} + n_{dk.}^{di})] \approx \log(\alpha_{dk} + \mathbb{E}_q[n_{dk.}^{di}]) - \frac{\text{var}_q[n_{dk.}^{di}]}{2(\alpha_{dk} + \mathbb{E}_q[n_{dk.}^{di}])^2}$$

$$\gamma_{dik} \propto \exp \left(\mathbb{E}_{q(C \setminus di)} \left[\log(\alpha_{dk} + n_{dk.}^{di}) + \log(\beta_{kw_{di}} + n_{.kw_{di}}^{di}) - \log \left(\sum_v \beta_{kv} + n_{.kv}^{di} \right) \right] \right)$$

$$\mathbb{E}_q[\log(\alpha_{dk} + n_{dk.}^{di})] \approx \log(\alpha_{dk} + \mathbb{E}_q[n_{dk.}^{di}]) - \frac{\text{var}_q[n_{dk.}^{di}]}{2(\alpha_{dk} + \mathbb{E}_q[n_{dk.}^{di}])^2}$$

$$\begin{aligned} \gamma_{dik} &\propto (\alpha_{dk} + \mathbb{E}[n_{dk.}^{di}])(\beta_{kw_{di}} + \mathbb{E}[n_{.kw_{di}}^{di}]) \left(\sum_v \beta_{kv} + \mathbb{E}[n_{.kv}^{di}] \right)^{-1} \\ &\quad \cdot \exp \left(-\frac{\text{var}_q[n_{dk.}^{di}]}{2(\alpha_{dk} + \mathbb{E}_q[n_{dk.}^{di}])^2} - \frac{\text{var}_q[n_{.kw_{di}}^{di}]}{2(\beta_{kw_{di}} + \mathbb{E}_q[n_{.kw_{di}}^{di}])^2} + \frac{\text{var}_q[n_{.k.}^{di}]}{2(\sum_v \beta_{kv} + \mathbb{E}_q[n_{.kv}^{di}])^2} \right) \end{aligned}$$

$$\gamma_{dik} \propto (\alpha_{dk} + \mathbb{E}[n_{dk.}^{di}]) (\beta_{kw_{di}} + \mathbb{E}[n_{.kw_{di}}^{di}]) \left(\sum_v \beta_{kv} + \mathbb{E}[n_{.kv}^{di}] \right)^{-1} \\ \cdot \exp \left(-\frac{\text{var}_q[n_{dk.}^{di}]}{2(\alpha_{dk} + \mathbb{E}_q[n_{dk.}^{di}])^2} - \frac{\text{var}_q[n_{.kw_{di}}^{di}]}{2(\beta_{kw_{di}} + \mathbb{E}_q[n_{.kw_{di}}^{di}])^2} + \frac{\text{var}_q[n_{.k.}^{di}]}{2(\sum_v \beta_{kv} + \mathbb{E}_q[n_{.kv}^{di}])^2} \right)$$

Note that γ_{dik} doesn't depend on $i \in 1, \dots, l_d$, it's the same for all w_{di} in d with $w_{di} = v$!

- ▶ memory requirement: $\mathcal{O}(DKV)$, since we have to store γ_{dik} for each value of $i \in 1, \dots, V$ and
 - ▶ $\mathbb{E}[n_{dk.}], \text{var}[n_{dk.}] \in \mathbb{R}^{D \times K}$
 - ▶ $\mathbb{E}[n_{.kv}], \text{var}[n_{.kv}] \in \mathbb{R}^{K \times V}$
 - ▶ $\mathbb{E}[n_{.k.}], \text{var}[n_{.k.}] \in \mathbb{R}^K$
- ▶ computational complexity: $\mathcal{O}(DKV)$ We can loop over V rather than l_d (good for long documents!) Often, a document will be sparse in V , so iteration cost can be much lower.

Designing a probabilistic machine learning method:

1. get the **data**
 - 1.1 try to collect as much meta-data as possible
2. build the **model**
 - 2.1 identify quantities and datastructures; assign names
 - 2.2 design a generative process (graphical model)
 - 2.3 assign (conditional) distributions to factors/arrows (use exponential families!)
3. design the **algorithm**
 - 3.1 consider conditional independence
 - 3.2 try standard methods for early experiments
 - 3.3 run unit-tests and sanity-checks
 - 3.4 identify bottlenecks, find customized approximations and refinements

Simple, generic algorithms or toolboxes can give an idea of a model's efficacy.
 To make an efficient product, you might have to build your own approximation, and tune it.

